

AS91946: Interpret and Apply Mathematical and Statistical Information

NCEA Level 1 Statistics — Year 11

Anthony R Davidson

July 2026

Table of contents

1	Welcome	1
	Welcome	2
1.1	What this standard is about	2
1.2	How to use this book	2
1.3	Tools we use	2
1.4	Getting started	3
2	What is Statistics?	4
3	What is Statistics?	5
3.1	Why statistics matters	5
3.2	The statistical enquiry cycle (PPDAC)	5
3.3	Types of statistical questions	6
3.4	Why context matters	7
3.5	Glossary	7
3.6	Check your understanding	7
4	Types of Data and Variables	8
5	Types of Data and Variables	9
5.1	Categorical vs numerical data	9
5.2	Categorical data	9
5.3	Numerical data	10
5.4	Real data example	10
5.5	Choosing the right display	11
5.6	Glossary	11
5.7	Check your understanding	11
6	Statistical Graphs and Displays	13
7	Statistical Graphs and Displays	14
7.1	Setting up our data	14
7.2	Bar charts	14
7.3	Dot plots	15
7.4	Histograms	16
7.5	Box plots (box-and-whisker plots)	18
7.6	Stem-and-leaf plots	19
7.7	Scatter plots	20
7.8	Pie charts	21
7.9	Choosing the right graph	22
7.10	Check your understanding	22
8	Measures of Centre and Spread	24

9 Measures of Centre and Spread	25
9.1 Measures of centre	25
9.2 Measures of spread	27
9.3 Comparing groups	29
9.4 Identifying outliers	30
9.5 Glossary	30
9.6 Check your understanding	31
10 Probability	32
11 Probability	33
11.1 Probability scale	33
11.2 Key vocabulary	33
11.3 Simple probability — equally likely outcomes	33
11.4 Complementary events	34
11.5 Experimental (relative frequency) probability	34
11.6 Two-way tables	36
11.7 Tree diagrams	37
11.8 Expected value (extension)	37
11.9 Glossary	38
11.10 Check your understanding	38
12 Interpreting Statistical Information	39
13 Interpreting Statistical Information	40
13.1 Reading statistical claims	40
13.2 Sources of bias	40
13.3 Sample size and reliability	40
13.4 Correlation and causation	42
13.5 Interpreting graphs critically	42
13.6 Reading media reports — a framework	44
13.7 A worked example	45
13.8 Glossary	45
13.9 Check your understanding	45
14 Practice Questions and Exam Tips	47
15 Practice Questions and Achievement Standard Tips	48
15.1 About the achievement standard	48
15.2 What each grade requires	48
15.3 Writing tips for the exam	49
15.4 Practice set 1 — Achieved level	49
15.5 Practice set 2 — Merit level	50
15.6 Practice set 3 — Excellence level	52
15.7 Key formulas summary	52
15.8 Exam strategy checklist	53
15.9 Useful resources	53
16 References	54
References	55
16.1 Achievement Standard	55

Table of contents

16.2 New Zealand Statistics Curriculum	55
16.3 Textbooks and Teaching Resources	55
16.4 Data Sources	55
16.5 Software	55
16.6 Statistical Literacy	56

List of Figures

1.1	An example of the kind of visualisations we will create in this book.	3
3.1	The PPDAC cycle — a framework for statistical investigations.	6
5.1	Classification of data types.	9
7.1	Bar chart showing how Year 11 students travel to school.	15
7.2	Dot plot of daily study hours for Year 11 students.	16
7.3	Histogram of student heights (cm).	17
7.4	Examples of different distribution shapes.	18
7.5	Box plot of maths scores by transport type.	19
7.6	Anatomy of a box plot.	19
7.7	Scatter plot of study hours vs maths score.	21
7.8	Pie chart of transport modes.	22
9.1	Effect of outliers on mean vs median.	27
9.2	Comparing maths scores between Group A and Group B.	29
11.1	The probability scale from impossible (0) to certain (1).	33
11.2	Relative frequency of heads approaches 0.5 as the number of flips increases (the law of large numbers).	35
11.3	Tree diagram for drawing two balls from a bag (with replacement).	37
13.1	Different types of bias and how they affect results.	41
13.2	Variability in sample estimates decreases as sample size increases.	42
13.3	A spurious correlation: ice cream sales and drownings both increase in summer.	43
13.4	The same data displayed with a truncated y-axis (left) vs starting at zero (right).	44
15.1	Weekly exercise hours for adults and teenagers.	51

List of Tables

5.1	Example student survey data — identify the type of each variable.	10
9.4	Summary statistics for Groups A and B	30
11.2	Sport preference by gender ($n = 120$)	36
15.3	Weekly exercise hours: adults vs teenagers	50

1 Welcome

Welcome

This resource has been developed to support Year 11 students preparing for **NCEA Level 1 Achievement Standard 91946**: *Interpret and apply mathematical and statistical information in context* (4 credits).

1.1 What this standard is about

In this standard you will learn to:

- **Read and interpret** statistical displays, tables, and graphs
- **Apply** mathematical and statistical ideas to real-world problems
- **Communicate** findings clearly using statistical language
- **Evaluate** statistical claims made in media and everyday life

1.2 How to use this book

Each chapter builds on the previous one. You will find:

- Clear explanations of key concepts with New Zealand examples
- Interactive R code blocks showing how statistics is done with real data
- Practice questions at Achievement, Merit, and Excellence levels
- Glossary boxes highlighting important vocabulary

Achievement Standard at a Glance

Grade	What you need to do
Achieved	Interpret and apply statistical information
Merit	Interpret and apply statistical information with justification
Excellence	Interpret and apply statistical information with statistical insight

1.3 Tools we use

All data visualisations and calculations in this book are produced using **R** and rendered with **Quarto**. You do not need to be able to write R code to succeed in this standard — the code blocks are provided so your teacher can demonstrate concepts live and so you can explore data yourself if you choose.




Figure 1.1: An example of the kind of visualisations we will create in this book.

1.4 Getting started

Head to Chapter [3](#) to begin your statistics journey!

2 What is Statistics?

3 What is Statistics?

3.1 Why statistics matters

Every day we are surrounded by data. News headlines cite poll results, sports commentators compare player averages, health campaigns report risk percentages, and social media shows us trending topics. **Statistics** is the science that helps us make sense of all this information.

i Key definition

Statistics is the practice of collecting, organising, summarising, analysing, and interpreting data in order to draw conclusions and make decisions.

3.2 The statistical enquiry cycle (PPDAC)

New Zealand's statistics curriculum uses the **PPDAC cycle** as a framework for carrying out a statistical investigation:

Stage	What you do
Problem	Pose a statistical question
Plan	Plan how to collect data
Data	Collect and record data
Analysis	Explore and summarise the data
Conclusion	Interpret results and answer the question

The PPDAC Cycle




Figure 3.1: The PPDAC cycle — a framework for statistical investigations.

3.3 Types of statistical questions

Not all questions are statistical questions. A statistical question expects **variability** in the answers.

Tip

Statistical question: “How tall are the students in Year 11?”

(Heights will vary — this is statistical!)

Not a statistical question: “How tall is Ms Smith?”

(Only one answer — not statistical.)

3.3.1 Practice identifying statistical questions

Which of the following are statistical questions?

1. How many pages are in Harry Potter and the Philosopher’s Stone?
2. How many hours per week do Year 11 students spend on homework?
3. What is the population of New Zealand?
4. What are the most popular music genres among teenagers in Auckland?
5. What colour is the school uniform?

Answers

1. Not statistical — there is one fixed answer.
2. Statistical — different students spend different amounts of time.
3. Not statistical — there is one (changing but definite) answer.
4. Statistical — different teenagers have different preferences.

5. Not statistical — the uniform colour is fixed.

3.4 Why context matters

In AS91946, **context** is central. You will always work with data that relates to real situations. When interpreting statistics, you must:

- **Name the variable** being measured
- **State the units** (e.g. centimetres, dollars, minutes)
- **Refer to the population** the data came from
- **Relate findings back** to the real-world situation

Common mistake

Saying “the mean is 24” is not enough. Say
“The mean age of survey respondents was 24 years.”

3.5 Glossary

Term	Meaning
Data	Values collected by observation or measurement
Variable	A characteristic that can take different values
Population	The entire group of interest
Sample	A subset of the population selected for study
Statistic	A numerical summary calculated from a sample
Parameter	A numerical summary of a whole population

3.6 Check your understanding

1. In your own words, explain what the PPDAC cycle is.
2. Give two examples of statistical questions you could investigate at your school.
3. A student says “the average score is 65”. What is missing from this statement?
4. What is the difference between a population and a sample?

4 Types of Data and Variables

5 Types of Data and Variables

Understanding the **type of data** you have is the first step in choosing the right graph or summary statistic. The wrong choice can make data hard to understand — or even misleading.

5.1 Categorical vs numerical data




Figure 5.1: Classification of data types.

5.2 Categorical data

Categorical (qualitative) data describes qualities or categories. The values are labels, not numbers.

5.2.1 Nominal data

Nominal data has **no natural order**. Categories are just names.

Examples: - Eye colour (brown, blue, green, grey) - Favourite sport (rugby, netball, football, basketball) - Type of transport to school (bus, car, walk, bike)

5.2.2 Ordinal data

Ordinal data has a **natural order**, but the gaps between categories are not equal.

Examples: - Survey ratings (strongly disagree → disagree → neutral → agree → strongly agree) - Year group (Year 9, Year 10, Year 11, Year 12, Year 13) - Clothing size (XS, S, M, L, XL)

5.3 Numerical data

Numerical (quantitative) data represents quantities that can be measured or counted.

5.3.1 Discrete data

Discrete data can only take **specific, separate values** (usually whole numbers from counting).

Examples: - Number of siblings - Number of text messages sent per day - Goals scored in a rugby match

5.3.2 Continuous data

Continuous data can take **any value within a range** (usually from measuring).

Examples: - Height (cm) - Weight (kg) - Time to complete a race (seconds) - Temperature (°C)

5.4 Real data example

```
# Simulated class survey data
set.seed(7)
n <- 10
survey <- data.frame(
  Student = paste0("S", 1:n),
  Gender = sample(c("Male", "Female", "Non-binary"), n, replace = TRUE),
  Height_cm = round(rnorm(n, 168, 10), 1),
  Siblings = sample(0:4, n, replace = TRUE),
  Transport = sample(c("Bus", "Car", "Walk", "Bike"), n, replace = TRUE),
  Rating = sample(c("Poor", "Fair", "Good", "Excellent"), n, replace = TRUE)
)

knitr::kable(survey, align = "c")
```

Table 5.1: Example student survey data — identify the type of each variable.

Student	Gender	Height_cm	Siblings	Transport	Rating
S1	Female	154.3	1	Bus	Fair
S2	Non-binary	144.2	3	Bike	Excellent
S3	Non-binary	163.2	2	Bus	Excellent
S4	Non-binary	162.6	3	Bus	Poor
S5	Female	181.2	0	Car	Excellent
S6	Non-binary	152.9	2	Bike	Fair

Student	Gender	Height_cm	Siblings	Transport	Rating
S7	Female	167.8	3	Bus	Good
S8	Female	164.5	0	Walk	Excellent
S9	Non-binary	161.7	2	Bus	Fair
S10	Female	159.0	4	Walk	Fair

💡 Identify the variable types

- **Gender** → Categorical, Nominal
- **Height_cm** → Numerical, Continuous
- **Siblings** → Numerical, Discrete
- **Transport** → Categorical, Nominal
- **Rating** → Categorical, Ordinal

5.5 Choosing the right display

The type of data determines the best graph to use:

Data type	Suitable displays
Categorical (nominal/ordinal)	Bar chart, pie chart, dot plot
Numerical (discrete)	Dot plot, bar chart, stem-and-leaf
Numerical (continuous)	Histogram, box plot, dot plot, stem-and-leaf
Two numerical variables	Scatter plot

5.6 Glossary

Term	Meaning
Variable	A characteristic that varies between individuals
Categorical	Data that falls into named groups or categories
Nominal	Categorical with no natural order
Ordinal	Categorical with a natural order
Numerical	Data that can be measured or counted
Discrete	Numerical data that can only take specific values
Continuous	Numerical data that can take any value in a range

5.7 Check your understanding

Classify each variable as **categorical (nominal/ordinal)** or **numerical (discrete/continuous)**:

1. Shoe size
2. Favourite subject at school
3. Time spent sleeping (hours per night)
4. Number of books read last month
5. Blood type
6. Temperature in a classroom

7. Year level
8. Distance from home to school (km)

i Answers

1. Numerical, discrete (or ordinal — shoe sizes have a natural order but are not continuous)
2. Categorical, nominal
3. Numerical, continuous
4. Numerical, discrete
5. Categorical, nominal
6. Numerical, continuous
7. Categorical, ordinal (or numerical discrete depending on context)
8. Numerical, continuous

6 Statistical Graphs and Displays

7 Statistical Graphs and Displays

A well-chosen graph makes patterns in data immediately visible. In this chapter you will learn to **create**, **read**, and **critically interpret** the most common statistical displays.

7.1 Setting up our data

Throughout this chapter we use a simulated survey of 60 Year 11 students across a New Zealand school. The data includes variables on transport, height, study time, and maths scores.

```
set.seed(2024)
n <- 60

students <- data.frame(
  transport = sample(c("Bus", "Car", "Walk", "Bike"), n,
                    replace = TRUE, prob = c(0.4, 0.25, 0.25, 0.1)),
  height_cm = round(rnorm(n, mean = 167, sd = 9), 1),
  study_hrs = round(pmax(0, rnorm(n, mean = 2.5, sd = 1.2)), 1),
  maths_score = round(pmin(100, pmax(0, rnorm(n, mean = 65, sd = 15))))
)
```

7.2 Bar charts

Bar charts display the **frequency** (count) or **relative frequency** (proportion) of categorical data.

```
transport_counts <- table(students$transport)

barplot(
  sort(transport_counts, decreasing = TRUE),
  col = c("#4dac26", "#b8e186", "#f1b6da", "#d7191c"),
  main = "How Year 11 Students Travel to School",
  xlab = "Mode of Transport",
  ylab = "Number of Students",
  border = "white",
  las = 1
)
```




Figure 7.1: Bar chart showing how Year 11 students travel to school.

7.2.1 Reading a bar chart

💡 Tip

When describing a bar chart, always comment on:

1. **The most and least common** categories
2. **The overall pattern** (are categories roughly equal, or is one dominant?)
3. **The context** — what does this tell us about the group?

Sample answer: “Bus is the most common way students travel to school, with over 20 students using this method. Bike is the least common. More students use public or active transport (bus, walk, bike) than passive transport (car).”

7.3 Dot plots

Dot plots are useful for **small numerical datasets**. Each data point is shown as a dot above a number line.

```
# Rounded to nearest 0.5 for clearer dot plot
study_rounded <- round(students$study_hrs * 2) / 2

stripchart(
  study_rounded,
  method = "stack",
  pch    = 19,
  col    = "#0571b0",
  cex    = 1.4,
```

```

xlab  = "Daily Study Hours",
main  = "Daily Study Hours - Year 11 Students",
offset = 0.4
)

```




Figure 7.2: Dot plot of daily study hours for Year 11 students.

7.4 Histograms

A histogram shows the **distribution of continuous numerical data** by grouping values into intervals (bins).

```

hist(
  students$height_cm,
  breaks = 10,
  col    = "#2c7bb6",
  border = "white",
  main   = "Distribution of Student Heights",
  xlab   = "Height (cm)",
  ylab   = "Frequency",
  las    = 1
)

```




Figure 7.3: Histogram of student heights (cm).

7.4.1 Describing a histogram — SOCS

Use **SOCS** to describe a distribution:

Letter	What to describe
Shape	Symmetric, skewed left, skewed right, uniform, bimodal
Outliers	Any unusual values far from the main group
Centre	Where the middle of the data is (mean or median)
Spread	How variable the data is (range, IQR, standard deviation)

```
par(mfrow = c(1, 3))

# Symmetric
set.seed(1)
hist(rnorm(500, 50, 10), breaks = 15, col = "#4dac26", border = "white",
     main = "Symmetric", xlab = "", ylab = "Frequency", las = 1)

# Right-skewed
hist(rexp(500, 0.1), breaks = 15, col = "#f4a582", border = "white",
     main = "Right (positive) skew", xlab = "", ylab = "", las = 1)

# Left-skewed
hist(100 - rexp(500, 0.1), breaks = 15, col = "#92c5de", border = "white",
     main = "Left (negative) skew", xlab = "", ylab = "", las = 1)

par(mfrow = c(1, 1))
```




Figure 7.4: Examples of different distribution shapes.

7.5 Box plots (box-and-whisker plots)

A box plot summarises a **five-number summary**: Minimum, Q1, Median, Q3, Maximum.

```
boxplot(
  maths_score ~ transport,
  data    = students,
  col     = c("#4dac26", "#f4a582", "#92c5de", "#d7191c"),
  border  = "#333333",
  main    = "Maths Scores by Transport Mode",
  xlab    = "Mode of Transport",
  ylab    = "Maths Score",
  las     = 1
)
```





Figure 7.5: Box plot of maths scores by transport type.

7.5.1 Parts of a box plot




Figure 7.6: Anatomy of a box plot.

7.6 Stem-and-leaf plots

Stem-and-leaf plots show the actual data values and the shape of the distribution simultaneously. They work well for **small datasets** (up to ~50 values).

```
# Show a stem-and-leaf for a subset of maths scores
cat("Stem-and-leaf plot of maths scores (first 30 students):\n\n")
#> Stem-and-leaf plot of maths scores (first 30 students):
stem(students$maths_score[1:30])
#>
#>   The decimal point is 1 digit(s) to the right of the |
#>
#>    2 | 1
#>    3 | 9
#>    4 | 8
#>    5 | 3468999
#>    6 | 04557789
#>    7 | 0169
#>    8 | 23448
#>    9 | 18
#>   10 | 0
```

7.7 Scatter plots

Scatter plots show the **relationship between two numerical variables**.

```
plot(
  students$study_hrs, students$maths_score,
  pch = 19,
  col = adjustcolor("#0571b0", alpha.f = 0.7),
  xlab = "Daily Study Hours",
  ylab = "Maths Score",
  main = "Study Hours vs Maths Score - Year 11 Students"
)
abline(lm(maths_score ~ study_hrs, data = students),
  col = "#d7191c", lwd = 2)
```




Figure 7.7: Scatter plot of study hours vs maths score.

7.7.1 Describing a scatter plot

Comment on:

1. **Direction** — positive, negative, or no association
2. **Form** — linear or non-linear
3. **Strength** — strong, moderate, or weak
4. **Outliers** — any unusual points

⚠ Correlation Causation

Just because two variables are related **does not** mean one causes the other. Always think about whether an association makes sense in context.

7.8 Pie charts

Pie charts show **proportions** of categorical data. Use them carefully — bar charts are usually easier to read precisely.

```
pie(
  transport_counts,
  col = c("#4dac26", "#f4a582", "#92c5de", "#d7191c"),
  main = "Transport Modes - Year 11 Students",
  labels = paste0(names(transport_counts), "\n(",
    round(100 * transport_counts / sum(transport_counts)), "%")
)
```

Transport Modes — Year 11 Students




Figure 7.8: Pie chart of transport modes.

7.9 Choosing the right graph

Situation	Best graph
One categorical variable	Bar chart, pie chart
One numerical variable (small n)	Dot plot, stem-and-leaf
One numerical variable (large n)	Histogram, box plot
Comparing groups numerically	Side-by-side box plots, back-to-back stem-and-leaf
Two numerical variables	Scatter plot

7.10 Check your understanding

1. A student recorded the heights of 25 classmates. Which graph type is most appropriate to display this data? Explain your choice.
2. Describe the scatter plot in Figure 7.7 using the four features (direction, form, strength, outliers).
3. A bar chart shows the following transport counts: Bus = 24, Car = 15, Walk = 15, Bike = 6. What percentage of students walk to school?
4. True or false: “A right-skewed distribution has a longer tail on the left.” Explain.

i Answers

1. A histogram or dot plot — height is numerical and continuous.
2. Positive direction (more study → higher score), roughly linear form, moderate strength, possibly one or two outliers.

3. Walk = $15/60 = 25\%$.
4. False — a right-skewed (positively skewed) distribution has a longer tail on the RIGHT.

8 Measures of Centre and Spread

9 Measures of Centre and Spread

Once you have a graph, you need **numerical summaries** to describe and compare distributions precisely. The two key things to describe are:

- **Centre** — where is the “middle” of the data?
- **Spread** — how variable is the data?

9.1 Measures of centre

9.1.1 The mean

The **mean** (average) is calculated by adding all values and dividing by the number of values:

$$\bar{x} = \frac{\sum x_i}{n}$$

```
set.seed(2024)
scores <- c(55, 62, 68, 70, 72, 75, 78, 80, 83, 85,
            88, 60, 91, 74, 66, 79, 52, 95, 73, 71)
n <- length(scores)

cat("Scores:", sort(scores), "\n")
#> Scores: 52 55 60 62 66 68 70 71 72 73 74 75 78 79 80 83 85 88 91 95
cat("Mean: ", round(mean(scores), 2), "\n")
#> Mean: 73.85
cat("n = ", n, "\n")
#> n = 20
```

9.1.2 The median

The **median** is the middle value when data is sorted from smallest to largest. It is not affected by outliers.

- If n is **odd**: median = middle value
- If n is **even**: median = average of the two middle values

```
sorted_scores <- sort(scores)
cat("Sorted scores:\n", sorted_scores, "\n\n")
#> Sorted scores:
#> 52 55 60 62 66 68 70 71 72 73 74 75 78 79 80 83 85 88 91 95
cat("n = 20 (even) → median = mean of 10th and 11th values\n")
#> n = 20 (even) → median = mean of 10th and 11th values
cat("10th value:", sorted_scores[10], "\n")
#> 10th value: 73
cat("11th value:", sorted_scores[11], "\n")
#> 11th value: 74
```

```
cat("Median =", median(scores), "\n")
#> Median = 73.5
```

9.1.3 The mode

The **mode** is the most frequently occurring value. Data can have: - **No mode** (all values appear once) - **One mode** (unimodal) - **Two modes** (bimodal) - **More than two modes** (multimodal)

```
# Custom mode function (base R has no built-in mode for numeric)
stat_mode <- function(x) {
  tbl <- table(x)
  as.numeric(names(tbl)[tbl == max(tbl)])
}

test_data <- c(4, 7, 2, 7, 9, 7, 3, 4, 5, 4)
cat("Data:", test_data, "\n")
#> Data: 4 7 2 7 9 7 3 4 5 4
cat("Mode(s):", stat_mode(test_data), "\n")
#> Mode(s): 4 7
```

9.1.4 When to use each measure of centre

Measure	Best used when...	Affected by outliers?
Mean	Data is roughly symmetric, no outliers	Yes
Median	Data is skewed or has outliers	No
Mode	Data is categorical; or finding most common value	No

```
par(mfrow = c(1, 2))

# Without outlier
x1 <- c(10, 12, 13, 14, 15, 16, 17, 18, 19, 20)
hist(x1, col = "#4dac26", border = "white", main = "No Outlier",
     xlab = "Value", ylab = "Frequency", xlim = c(0, 80), breaks = 5)
abline(v = mean(x1), col = "#d7191c", lwd = 2, lty = 1)
abline(v = median(x1), col = "#0571b0", lwd = 2, lty = 2)
legend("topright", c(paste("Mean =", mean(x1)),
                      paste("Median =", median(x1))),
     col = c("#d7191c", "#0571b0"), lwd = 2, lty = c(1, 2), bty = "n")

# With outlier
x2 <- c(x1, 75)
hist(x2, col = "#f4a582", border = "white", main = "With Outlier (75)",
     xlab = "Value", ylab = "", xlim = c(0, 80), breaks = 8)
abline(v = mean(x2), col = "#d7191c", lwd = 2, lty = 1)
abline(v = median(x2), col = "#0571b0", lwd = 2, lty = 2)
legend("topright", c(paste("Mean =", round(mean(x2), 1)),
                      paste("Median =", median(x2))),
```



```
col = c("#d7191c", "#0571b0"), lwd = 2, lty = c(1, 2), bty = "n")
par(mfrow = c(1, 1))
```




Figure 9.1: Effect of outliers on mean vs median.

9.2 Measures of spread

9.2.1 Range

The **range** is the simplest measure of spread:

$$\text{Range} = \text{Maximum} - \text{Minimum}$$

```
cat("Range of scores:", max(scores) - min(scores), "\n")
#> Range of scores: 43
cat(" (Max =", max(scores), ", Min =", min(scores), ") \n")
#> (Max = 95 , Min = 52 )
```

9.2.2 Quartiles and the interquartile range (IQR)

Quartiles divide sorted data into four equal groups:

Quartile	Description
Q1 (lower quartile)	25% of data falls below this value
Q2 (median)	50% of data falls below this value
Q3 (upper quartile)	75% of data falls below this value

The **interquartile range (IQR)** is the spread of the middle 50% of data:

$$\text{IQR} = Q3 - Q1$$

```
q <- quantile(scores)
cat("Five-number summary:\n")
#> Five-number summary:
cat("  Minimum:", q[1], "\n")
#>   Minimum: 52
cat("   Q1:    ", q[2], "\n")
#>    Q1:     67.5
cat("  Median: ", q[3], "\n")
#>   Median: 73.5
cat("   Q3:    ", q[4], "\n")
#>    Q3:     80.75
cat("  Maximum:", q[5], "\n")
#>   Maximum: 95
cat("   IQR =", q[4] - q[2], "\n")
#>    IQR = 13.25
```

9.2.3 Standard deviation

The **standard deviation** measures how spread out data values are from the mean. A larger standard deviation means more variability.

$$s = \sqrt{\frac{\sum (x_i - \bar{x})^2}{n - 1}}$$

```
cat("Mean:", round(mean(scores), 2), "\n")
#> Mean: 73.85
cat("Standard deviation:", round(sd(scores), 2), "\n")
#> Standard deviation: 11.52
```

Rule of thumb

For roughly symmetric distributions, about **68%** of data falls within one standard deviation of the mean, and **95%** falls within two standard deviations (the empirical rule / 68-95-99.7 rule).

9.2.4 Choosing the right measure of spread

Measure	Use with...	Advantage
Range	Quick overview	Easy to calculate
IQR	Skewed data or when outliers present	Resistant to outliers
Standard deviation	Symmetric data, no outliers	Uses all data values

9.3 Comparing groups

A common task in AS91946 is to **compare two or more groups**. Always compare both centre AND spread.

```
set.seed(42)
group_a <- round(rnorm(30, mean = 65, sd = 10))
group_b <- round(rnorm(30, mean = 72, sd = 15))

boxplot(
  list("Group A" = group_a, "Group B" = group_b),
  col    = c("#4dac26", "#0571b0"),
  border = "#333333",
  main   = "Maths Scores: Group A vs Group B",
  ylab   = "Score",
  las    = 1
)
```




Figure 9.2: Comparing maths scores between Group A and Group B.

```
summary_stats <- data.frame(
  Group    = c("Group A", "Group B"),
  n        = c(length(group_a), length(group_b)),
  Mean     = round(c(mean(group_a), mean(group_b)), 1),
  Median   = c(median(group_a), median(group_b)),
  SD       = round(c(sd(group_a), sd(group_b)), 1),
  IQR      = c(IQR(group_a), IQR(group_b)),
  Min      = c(min(group_a), min(group_b)),
  Max      = c(max(group_a), max(group_b))
)
```

```
knitr::kable(summary_stats, caption = "Summary statistics for Groups A and B")
```

Table 9.4: Summary statistics for Groups A and B

Group	n	Mean	Median	SD	IQR	Min	Max
Group A	30	65.6	64	12.6	14.25	38	88
Group B	30	70.2	74	15.8	21.25	27	96

9.3.1 Model comparison paragraph (Merit level)

“Group B has a higher median score (72) than Group A (65), suggesting that Group B performed better overall. However, Group B also has a larger IQR (approximately 20 vs 14) and standard deviation (15.1 vs 9.8), indicating their scores were more variable. Group A’s performance was more consistent. Both groups appear to have roughly symmetric distributions with no obvious outliers.”

9.4 Identifying outliers

A value is often flagged as an **outlier** if it falls more than $1.5 \times \text{IQR}$ below Q1 or above Q3:

$$\text{Lower fence} = Q1 - 1.5 \times \text{IQR}$$

$$\text{Upper fence} = Q3 + 1.5 \times \text{IQR}$$

```
Q1 <- quantile(scores, 0.25)
Q3 <- quantile(scores, 0.75)
iqr <- Q3 - Q1

lower_fence <- Q1 - 1.5 * iqr
upper_fence <- Q3 + 1.5 * iqr

cat("Q1 =", Q1, "\n")
#> Q1 = 67.5
cat("Q3 =", Q3, "\n")
#> Q3 = 80.75
cat("IQR =", iqr, "\n")
#> IQR = 13.25
cat("Lower fence:", lower_fence, "\n")
#> Lower fence: 47.625
cat("Upper fence:", upper_fence, "\n")
#> Upper fence: 100.625
cat("Outliers:", scores[scores < lower_fence | scores > upper_fence], "\n")
#> Outliers:
```

9.5 Glossary

Term	Meaning
Mean	Sum of all values divided by the count
Median	Middle value when data is sorted
Mode	Most frequently occurring value
Range	Maximum – Minimum
IQR	$Q3 - Q1$; spread of the middle 50%
Standard deviation	Average distance of values from the mean
Outlier	A value that lies far from the rest of the data
Five-number summary	Min, $Q1$, Median, $Q3$, Max

9.6 Check your understanding

Use the dataset: **12, 15, 18, 18, 20, 22, 25, 25, 28, 100**

1. Calculate the mean and median. Which is more appropriate as a measure of centre for this data? Why?
2. Find $Q1$, $Q3$, and the IQR.
3. Calculate the range.
4. Is 100 an outlier? Use the $1.5 \times \text{IQR}$ rule to justify your answer.
5. Sketch a box plot of this data.

i Answers

1. Mean = $(12+15+18+18+20+22+25+25+28+100)/10 = 283/10 = 28.3$; Median = $(20+22)/2 = 21$. The **median** is more appropriate because of the outlier (100) which pulls the mean up.
2. $Q1 = 18$, $Q3 = 25$, IQR = 7.
3. Range = $100 - 12 = 88$.
4. Upper fence = $25 + 1.5 \times 7 = 25 + 10.5 = 35.5$. Since $100 > 35.5$, **yes**, 100 is an outlier.
5. Sketch: whiskers from 12 to 35.5, box from 18 to 25, median line at 21, and a dot at 100 representing the outlier.

10 Probability

11 Probability

Probability is a measure of how likely an event is to occur. It is expressed as a number between 0 (impossible) and 1 (certain), or as a percentage.

$$P(\text{event}) = \frac{\text{number of favourable outcomes}}{\text{total number of possible outcomes}}$$

11.1 Probability scale




Figure 11.1: The probability scale from impossible (0) to certain (1).

11.2 Key vocabulary

Term	Definition
Experiment	A process with uncertain outcomes
Outcome	A single possible result
Event	One or more outcomes grouped together
Sample space	The set of all possible outcomes
Complementary event	All outcomes that are NOT in the event

11.3 Simple probability — equally likely outcomes

When all outcomes are **equally likely**, we can use the classical definition.

11.3.1 Example: rolling a die

```

sample_space <- 1:6
cat("Sample space:", sample_space, "\n")
#> Sample space: 1 2 3 4 5 6

# P(rolling a 4)
p_four <- sum(sample_space == 4) / length(sample_space)
cat("P(rolling a 4) =", p_four, "\n")
#> P(rolling a 4) = 0.1666667

# P(rolling an even number)
p_even <- sum(sample_space %% 2 == 0) / length(sample_space)
cat("P(even number) =", p_even, "\n")
#> P(even number) = 0.5

# P(rolling > 4)
p_gt4 <- sum(sample_space > 4) / length(sample_space)
cat("P(> 4) =", p_gt4, "\n")
#> P(> 4) = 0.3333333

```

11.4 Complementary events

The **complement** of event A (written A' or $\bar{A}$) contains all outcomes NOT in A.

$$P(A') = 1 - P(A)$$

Example: If the probability of rain tomorrow is 0.3, then:

$$P(\text{no rain}) = 1 - 0.3 = 0.7$$

```

p_rain <- 0.3
p_no_rain <- 1 - p_rain
cat("P(rain) =", p_rain, "\n")
#> P(rain) = 0.3
cat("P(no rain) = 1 -", p_rain, "=", p_no_rain, "\n")
#> P(no rain) = 1 - 0.3 = 0.7

```

11.5 Experimental (relative frequency) probability

When we cannot calculate probability theoretically, we **estimate** it by repeating an experiment many times:

$$P(\text{event}) \approx \frac{\text{number of times event occurred}}{\text{total number of trials}}$$

11.5.1 Simulating coin flips


```

set.seed(99)
flips <- sample(c("Heads", "Tails"), 1000, replace = TRUE)
cumulative_heads <- cumsum(flips == "Heads") / seq_along(flips)

plot(
  cumulative_heads,
  type = "l",
  col = "#0571b0",
  lwd = 1.5,
  xlab = "Number of Flips",
  ylab = "Relative Frequency of Heads",
  main = "Simulated Coin Flips",
  ylim = c(0, 1),
  las = 1
)
abline(h = 0.5, col = "#d7191c", lwd = 2, lty = 2)
legend("topright", c("Relative frequency", "True probability (0.5)"),
      col = c("#0571b0", "#d7191c"), lwd = 2, lty = c(1, 2), bty = "n")

```




Figure 11.2: Relative frequency of heads approaches 0.5 as the number of flips increases (the law of large numbers).

i Law of Large Numbers

As the number of trials increases, the **relative frequency** (experimental probability) gets closer and closer to the **true probability**. This is why surveys and experiments need large samples.

11.6 Two-way tables

A **two-way table** (contingency table) organises data by two categorical variables and is very useful for calculating probabilities.

```
# Simulated data: 120 students, sport preference by gender
set.seed(7)
n_students <- 120
sport_data <- data.frame(
  Gender = sample(c("Male", "Female"), n_students, replace = TRUE,
    prob = c(0.5, 0.5)),
  Sport = sample(c("Rugby", "Netball", "Football", "Basketball"),
    n_students, replace = TRUE, prob = c(0.3, 0.25, 0.25, 0.2))
)

tbl <- table(sport_data$Sport, sport_data$Gender)
tbl_with_total <- addmargins(tbl)
knitr::kable(tbl_with_total, caption = "Sport preference by gender (n = 120)")
```

Table 11.2: Sport preference by gender (n = 120)

	Female	Male	Sum
Basketball	16	11	27
Football	13	17	30
Netball	18	17	35
Rugby	16	12	28
Sum	63	57	120

11.6.1 Reading probabilities from a two-way table

Using the table above:

```
total <- sum(tbl)

cat("P(Male) =", sum(tbl[, "Male"]), "/", total, "=",
  round(sum(tbl[, "Male"]) / total, 3), "\n")
#> P(Male) = 57 / 120 = 0.475

cat("P(Rugby) =", sum(tbl["Rugby", ]), "/", total, "=",
  round(sum(tbl["Rugby", ]) / total, 3), "\n")
#> P(Rugby) = 28 / 120 = 0.233

cat("P(Male AND Rugby) =", tbl["Rugby", "Male"], "/", total, "=",
  round(tbl["Rugby", "Male"] / total, 3), "\n")
#> P(Male AND Rugby) = 12 / 120 = 0.1

cat("P(Rugby | Male) =", tbl["Rugby", "Male"], "/",
  sum(tbl[, "Male"]), "=",
  round(tbl["Rugby", "Male"] / sum(tbl[, "Male"]), 3), "\n")
#> P(Rugby | Male) = 12 / 57 = 0.211
```

The last probability — $P(\text{Rugby} \mid \text{Male})$ — is a **conditional probability**: the probability of preferring rugby, **given** the student is male.

11.7 Tree diagrams

Tree diagrams help calculate probabilities for **sequential events**.

Tree Diagram: Drawing 2 Balls (Red/Blue, with replacement)




Figure 11.3: Tree diagram for drawing two balls from a bag (with replacement).

Key rule: Multiply probabilities along each branch, add probabilities for different paths leading to the same outcome.

11.8 Expected value (extension)

For a simple scenario, the **expected value** is the long-run average outcome:

$$E(X) = \sum x \cdot P(x)$$

Example: A raffle has 200 tickets at \$2 each. First prize is \$100, second prize is \$50. What is the expected gain per ticket?

```

outcomes <- c(100, 50, 0)      # winnings
probs     <- c(1/200, 1/200, 198/200)
cost      <- 2

expected_win <- sum(outcomes * probs)
expected_gain <- expected_win - cost

cat("Expected winnings:", expected_win, "\n")
#> Expected winnings: 0.75
  
```

```

cat("Ticket cost:      ", cost, "\n")
#> Ticket cost:      2
cat("Expected net gain:", expected_gain, "\n")
#> Expected net gain: -1.25
cat("(A negative expected gain means this is not a good investment!)\n")
#> (A negative expected gain means this is not a good investment!)

```

11.9 Glossary

Term	Definition
Probability	Likelihood of an event, from 0 to 1
Sample space	All possible outcomes
Event	One or more outcomes
Complement	All outcomes NOT in the event
Relative frequency	Experimental estimate of probability
Conditional probability	$P(A \text{ given } B)$: probability of A given B occurred
Two-way table	Table showing two categorical variables together
Tree diagram	Diagram showing sequential outcomes and probabilities

11.10 Check your understanding

- A bag contains 4 red, 3 blue, and 3 yellow marbles. A marble is drawn at random. Find:
 - $P(\text{red})$
 - $P(\text{not red})$
 - $P(\text{blue or yellow})$
- A class of 30 students was surveyed: 18 play sport, 12 play an instrument, and 6 do both.
 - Draw a Venn diagram.
 - Find $P(\text{plays sport but not instrument})$.
 - Find $P(\text{neither sport nor instrument})$.
- If $P(A) = 0.4$ and $P(B) = 0.3$ and A and B are independent, find $P(A \text{ and } B)$.

i Answers

- $4/10 = 0.4$; b. $6/10 = 0.6$; c. $6/10 = 0.6$
- $12/30 = 0.4$; c. $(30 - 18 - 12 + 6)/30 = 6/30 = 0.2$
- $P(A \text{ and } B) = 0.4 \times 0.3 = 0.12$ (for independent events)

12 Interpreting Statistical Information

13 Interpreting Statistical Information

A major focus of AS91946 is the ability to **critically read and interpret** statistical information from a variety of sources — including newspapers, websites, government reports, and scientific studies.

13.1 Reading statistical claims

Statistical claims appear everywhere in everyday life. As a critical consumer of information, you need to ask:

1. **Who** collected the data and why?
2. **How** was the data collected (survey, experiment, census)?
3. **How many** people or things were measured (sample size)?
4. **Who** was sampled — does the sample represent the population?
5. **What** do the numbers actually mean in context?

13.2 Sources of bias

Bias is a systematic error that causes results to be unrepresentative.

Type of bias	Description	Example
Sampling bias	Sample does not represent population	Surveying only Year 11s about a school-wide issue
Response bias	People don't answer truthfully	"Do you cheat on tests?"
Non-response bias	Certain groups are less likely to respond	Online survey misses those without internet
Leading question bias	Question wording influences the answer	"Don't you agree that homework is too much?"
Volunteer bias	Only motivated people respond	Self-selected online polls

13.3 Sample size and reliability

Larger samples generally give more **reliable** results — but size alone is not enough if the sample is biased.

```
set.seed(42)
true_prop <- 0.55 # True proportion in population
sample_sizes <- c(10, 30, 50, 100, 200, 500, 1000)
```




Figure 13.1: Different types of bias and how they affect results.

```
n_simulations <- 1000

set.seed(321)
variability <- sapply(sample_sizes, function(n) {
  estimates <- replicate(n_simulations, mean(sample(c(TRUE, FALSE), n,
                                                    replace = TRUE,
                                                    prob = c(true_prop,
                                                            1 - true_prop))))
  sd(estimates)
})

plot(
  sample_sizes, variability,
  type = "b",
  pch = 19,
  col = "#0571b0",
  lwd = 2,
  xlab = "Sample Size (n)",
  ylab = "Standard Deviation of Estimates",
  main = "Sample Size vs Variability of Estimates",
  las = 1
)
abline(h = 0, col = "#d7191c", lty = 2)
```




Figure 13.2: Variability in sample estimates decreases as sample size increases.

13.4 Correlation and causation

One of the most important concepts in statistical literacy is understanding the difference between **correlation** and **causation**.

⚠ Correlation Causation

Two variables can be **correlated** (move together) without one causing the other. There may be a **confounding variable** responsible for both.

13.4.1 Famous spurious correlations

- Ice cream sales and shark attacks both increase in summer → they are both caused by **hot weather**, not each other.
- Countries with more TVs have longer life expectancy → both are driven by **wealth**, not TVs causing people to live longer.

13.5 Interpreting graphs critically

When reading statistical graphs, watch out for:

13.5.1 Misleading y-axis

A y-axis that doesn't start at zero can make small differences look huge.

```
par(mfrow = c(1, 2))
```





Figure 13.3: A spurious correlation: ice cream sales and drownings both increase in summer.

```
years <- 2020:2024
revenue <- c(980, 990, 995, 1005, 1010)

# Misleading
plot(years, revenue, type = "b", pch = 19, col = "#d7191c", lwd = 2,
      ylim = c(975, 1015), xlab = "Year", ylab = "Revenue ($M)",
      main = "Misleading: Truncated Axis", las = 1)

# Honest
plot(years, revenue, type = "b", pch = 19, col = "#4dac26", lwd = 2,
      ylim = c(0, 1100), xlab = "Year", ylab = "Revenue ($M)",
      main = "Honest: Starting at Zero", las = 1)

par(mfrow = c(1, 1))
```




Figure 13.4: The same data displayed with a truncated y-axis (left) vs starting at zero (right).

13.5.2 Percentage vs absolute numbers

💡 Be careful with percentages

“Crime increased by 100%” sounds alarming — but if there were only 2 crimes last year and 4 this year, the absolute increase is tiny. Always ask: **100% of what?**

13.5.3 Averages without context

The **mean** can be pulled by extreme values, making it misleading if used alone. Always ask what **measure of centre** was used and whether it is appropriate.

13.6 Reading media reports — a framework

Use these questions when evaluating a statistical claim in the media:

1. **What is being measured?** (the variable and its units)
2. **Who was studied?** (the population and sample)
3. **How was data collected?** (method and potential biases)
4. **How large was the sample?**
5. **What is the source?** (government, independent researcher, company with a vested interest?)
6. **Are conclusions supported by the evidence?** Or is the claim overstated?
7. **Are comparisons fair?** Are groups being compared actually comparable?

13.7 A worked example

i News headline

“Study shows eating chocolate improves exam scores by 15%”

Critical analysis:

- **What was measured?** Exam scores before and after eating chocolate.
- **Who was studied?** Not stated — red flag!
- **Sample size?** Not given — could be only 10 students.
- **Confounding variables?** Students who eat chocolate may be from wealthier families with better study resources.
- **Causation or correlation?** This is a correlational study; chocolate doesn’t necessarily *cause* better scores.
- **Who funded it?** If a chocolate company funded it, there is a conflict of interest.

Conclusion: The headline is likely exaggerated or misleading. We would need much more information before accepting this claim.

13.8 Glossary

Term	Definition
Bias	Systematic error that distorts results
Confounding variable	A hidden variable that explains an apparent relationship
Correlation	A statistical relationship between two variables
Causation	One variable directly causes a change in another
Reliable	Consistent results if repeated
Valid	Measuring what it claims to measure
Population	The full group of interest
Sample	A subset of the population

13.9 Check your understanding

1. A newspaper article claims: “People who eat breakfast earn more money than those who skip it.” Identify a possible confounding variable.
2. A survey found that 80% of students support a longer lunch break. The survey was conducted by standing outside the school canteen. What type of bias is this? How might it affect results?
3. A graph shows the number of electric vehicles sold in New Zealand, with the y-axis starting at 45,000 rather than 0. Sales increased from 47,000 to 51,000. On the graph, this looks like a 400% increase. Calculate the true percentage increase.
4. Explain the difference between correlation and causation using your own example.

i Answers

1. Wealthier people are more likely to eat breakfast (they can afford food) and also earn more — wealth is the confounding variable.
2. Sampling bias / volunteer bias — students near the canteen are more likely to want a longer lunch break. Results will overestimate support.
3. True % increase = $(51,000 - 47,000) / 47,000 \times 100 = 8.5\%$.
4. Accept any reasonable example, e.g., “Countries with more hospitals have higher death rates — hospitals don’t cause death; both are linked to population size.”

14 Practice Questions and Exam Tips

15 Practice Questions and Achievement Standard Tips

This chapter brings everything together with practice questions at each grade level, and advice on how to approach the AS91946 assessment.

15.1 About the achievement standard

Detail	Information
Standard number	AS91946
Title	Interpret and apply mathematical and statistical information in context
Level	NCEA Level 1
Credits	4
Assessment	External (end-of-year exam)
Strands	Statistics and Probability

15.2 What each grade requires

15.2.1 Achieved

- Correctly read values from graphs and tables
- Calculate basic statistics (mean, median, mode, range)
- State probabilities from given information
- Identify variable types
- Make simple statements about data in context

15.2.2 Merit

- **Justify** your answers with reasons
- Compare and contrast groups using both centre AND spread
- Identify potential sources of bias
- Explain the difference between correlation and causation
- Interpret probabilities and relate them to real-world context

15.2.3 Excellence

- Show **statistical insight** — go beyond the obvious
- Consider alternative explanations or limitations
- Make connections between different parts of the data
- Critically evaluate statistical claims
- Use appropriate statistical language precisely

15.3 Writing tips for the exam

The golden rule

Always refer back to the context. Every answer should mention: - The **variable** being described - The **units** of measurement - The **group** being described (population/sample)

“The mean is 65.”

“The mean maths score for Year 11 students was 65 out of 100.”

15.3.1 Comparison sentence starters

Purpose	Sentence starters
Centre	“On average...”; “The typical...”; “The median...is higher/lower than...”
Spread	“The scores are more/less variable...”; “The IQR is larger, suggesting...”
Shape	“The distribution is roughly symmetric / right-skewed...”
Outliers	“There appears to be an outlier at...which may indicate...”
Causation	“However, this may be due to...rather than a direct causal relationship”

15.4 Practice set 1 — Achieved level

#> Heights (cm) of 20 Year 11 students:

#> 147.3 153.9 154.5 157.3 158.9 159.5 163.6 164.3 164.8 165 165.2 166.4 166.5 167.2 169.3 170.1

Questions:

1. What type of data is height? Give a reason.
2. Calculate the mean height. Round to one decimal place.
3. Find the median height.
4. Find the range.
5. A student drew a histogram of this data. Circle the most appropriate description: **sym-metric** / **right-skewed** / **left-skewed**

i Worked solutions

```

set.seed(111)
heights <- round(rnorm(20, 168, 9), 1)

cat("1. Numerical, continuous - height can take any value within a range.\n\n")
#> 1. Numerical, continuous - height can take any value within a range.
cat("2. Mean:", round(mean(heights), 1), "cm\n\n")
#> 2. Mean: 164.8 cm
cat("3. Median:", median(heights), "cm\n\n")
#> 3. Median: 165.1 cm
cat("4. Range:", max(heights) - min(heights), "cm\n\n")
#> 4. Range: 37.3 cm
cat("5. Likely symmetric (heights tend to be normally distributed).\n")
#> 5. Likely symmetric (heights tend to be normally distributed).

```

15.5 Practice set 2 — Merit level

A local council surveyed 80 residents about their weekly exercise hours. The results were:

```

set.seed(222)
adults <- round(pmax(0, rnorm(40, mean = 4.5, sd = 2.2)), 1)
teenagers <- round(pmax(0, rnorm(40, mean = 6.2, sd = 3.1)), 1)

summary_tbl <- data.frame(
  Group = c("Adults", "Teenagers"),
  n = 40,
  Mean = round(c(mean(adults), mean(teenagers)), 2),
  Median = c(median(adults), median(teenagers)),
  SD = round(c(sd(adults), sd(teenagers)), 2),
  IQR = round(c(IQR(adults), IQR(teenagers)), 2),
  Min = c(min(adults), min(teenagers)),
  Max = c(max(adults), max(teenagers))
)

knitr::kable(summary_tbl, caption = "Weekly exercise hours: adults vs teenagers")

```

Table 15.3: Weekly exercise hours: adults vs teenagers

Group	n	Mean	Median	SD	IQR	Min	Max
Adults	40	4.45	4.35	2.07	2.85	0	7.9
Teenagers	40	6.58	7.10	2.85	3.58	0	13.0

```

boxplot(
  list(Adults = adults, Teenagers = teenagers),
  col = c("#4dac26", "#0571b0"),
  border = "#333333",

```



```
main = "Weekly Exercise Hours",
ylab = "Hours per week",
las = 1
)
```




Figure 15.1: Weekly exercise hours for adults and teenagers.

Questions:

1. Compare the centre of the two distributions. What does this tell us about exercise habits in context?
2. Compare the spread. What does this tell us?
3. The council used a telephone survey conducted at 2 pm on weekdays. Identify one source of bias and explain how it might affect the results.
4. Can we conclude from this data that teenagers exercise more because they are younger? Explain.

i Model answers

- 1. Centre:** The median exercise hours for teenagers (approx. 7.1 hrs) is higher than for adults (4.35 hrs). This suggests teenagers in this sample tend to do more weekly exercise than adults.
- 2. Spread:** The IQR and standard deviation are larger for teenagers (3.6 hrs vs 2.9 hrs). This means teenagers' exercise habits are more variable — some exercise a lot while others exercise very little. Adults' exercise habits are more consistent.
- 3. Bias:** Calling at 2 pm on weekdays means working adults are unlikely to answer (they are at work). This leads to **non-response bias** — the adult sample may over-represent retired or part-time workers, who may exercise more than the general adult population.
- 4. Causation:** No — we cannot conclude that age *causes* the difference in exercise. Teenagers may exercise more because they participate in school sport, not because of their age per se. Age and exercise are *correlated* in this sample, but there may be confounding

variables such as school PE requirements, free time, or lifestyle factors.

15.6 Practice set 3 — Excellence level

A researcher claims: “A study of 50 Auckland households found that households with more bookshelves have higher household income. Therefore, buying more bookshelves will increase your income.”

Questions:

1. Identify the two variables and their types.
2. Explain why this conclusion is flawed. Use statistical language.
3. Suggest one confounding variable that could explain the relationship.
4. What would the researcher need to do to properly establish causation?
5. The sample was 50 Auckland households selected by going door-to-door in Remuera (a wealthy suburb). Evaluate the reliability of generalising these findings to all New Zealand households.

i Model answers

1. Number of bookshelves: numerical, discrete. Household income: numerical, continuous.
2. The conclusion confuses **correlation with causation**. Even if the two variables are positively correlated, this does not mean bookshelves cause higher income. The relationship could be due to a third variable, and correlation studies cannot establish cause-and-effect without experimental control.
3. **Wealth/education level** — wealthier, more educated households are more likely to own books/bookshelves AND to have higher income. Both variables are driven by a common factor.
4. To establish causation, the researcher would need a **controlled experiment** — randomly assigning households to receive bookshelves or not, holding all other variables constant, and measuring income change over time. This would be impractical.
5. The sample has significant **sampling bias**. Remuera is one of Auckland’s wealthiest suburbs. The sample does not represent the diversity of New Zealand households in terms of income, region, or ethnicity. Results cannot be reliably generalised to all New Zealand households. A **stratified random sample** from across New Zealand would be needed.

15.7 Key formulas summary

Measure	Formula
Mean	$\bar{x} = \Sigma x / n$
Range	Max – Min
IQR	Q3 – Q1
Probability	$P(A) = \text{favourable outcomes} / \text{total outcomes}$
Complement	$P(A') = 1 - P(A)$
Lower outlier fence	$Q1 - 1.5 \times \text{IQR}$
Upper outlier fence	$Q3 + 1.5 \times \text{IQR}$

15.8 Exam strategy checklist

Before the exam, make sure you can:

- ☐ Identify variable types (categorical/numerical, nominal/ordinal, discrete/continuous)
- ☐ Read values accurately from bar charts, histograms, box plots, and scatter plots
- ☐ Calculate mean, median, mode, range, IQR
- ☐ Describe distributions using SOCS (Shape, Outliers, Centre, Spread)
- ☐ Compare two groups using both centre and spread in context
- ☐ Calculate simple probabilities and use complementary events
- ☐ Read and calculate from two-way tables
- ☐ Identify bias in data collection methods
- ☐ Distinguish correlation from causation
- ☐ Critically evaluate statistical claims in media

15.9 Useful resources

Resource	Description	URL
NZQA	Official achievement standard information	nzqa.govt.nz
CensusAtSchool	Real NZ student data for practice	new.censusatschool.org.nz
iNZight	Free NZ-made statistics software	inzight.nz
Stats NZ	Official NZ statistics and data	stats.govt.nz
Khan Academy	Free video lessons on statistics	khanacademy.org

16 References

References

16.1 Achievement Standard

New Zealand Qualifications Authority (NZQA). (2024). *Achievement Standard 91946: Interpret and apply mathematical and statistical information in context*. NZQA. <https://www.nzqa.govt.nz/nqfdocs/ncea-resource/achievements/2024/as91946.pdf>

16.2 New Zealand Statistics Curriculum

Ministry of Education New Zealand. (2007). *The New Zealand Curriculum*. Learning Media. <https://nzcurriculum.tki.org.nz>

Ministry of Education New Zealand. (2023). *Mathematics and Statistics — Achievement Objectives*. <https://nzcurriculum.tki.org.nz/The-New-Zealand-Curriculum/Mathematics-and-statistics>

16.3 Textbooks and Teaching Resources

Wild, C. J., & Seber, G. A. F. (2000). *Chance Encounters: A First Course in Data Analysis and Inference*. Wiley.

Pfannkuch, M., & Wild, C. J. (2004). Towards an understanding of statistical thinking. In D. Ben-Zvi & J. Garfield (Eds.), *The Challenge of Developing Statistical Literacy, Reasoning and Thinking* (pp. 17–46). Kluwer Academic.

16.4 Data Sources

Statistics New Zealand (Stats NZ). (2024). *Household Economic Survey*. <https://www.stats.govt.nz>

CensusAtSchool New Zealand. (2024). *CensusAtSchool data*. University of Auckland. <https://new.censusatschool.org.nz>

16.5 Software

R Core Team. (2024). *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing. <https://www.r-project.org>

Allaire, J. J., Teague, C., Scheidegger, C., Xie, Y., & Dervieux, C. (2024). *Quarto* (Version 1.5). <https://quarto.org>

Wickham, H., et al. (2023). *ggplot2: Elegant Graphics for Data Analysis* (3rd ed.). Springer. <https://ggplot2-book.org>

16.6 Statistical Literacy

Huff, D. (1954). *How to Lie with Statistics*. W. W. Norton & Company.

Wheelan, C. (2013). *Naked Statistics: Stripping the Dread from the Data*. W. W. Norton & Company.